

Monitoring the HPC cluster

Why monitor the cluster

When you submit a job you request a resource allocation in the cluster (number of nodes, number of CPUs, memory etc). When the job is executed all of those resources are allocated to that job; even if they're not being used as intended by your script they won't be available to anyone else whilst your job is running.

Resource availability

Checking the cluster utilisation before you submit a job allows you to adjust your script according to the resources available. Requesting too many resources for a single job may leave your job sitting in the queue for a long time, potentially days. Your job may get executed earlier if you request fewer resources (12 CPU's instead of 24 CPU's). Even if the job takes twice as long to compute it may still finish earlier if it progresses through the queue more quickly!

Code performance

You may be using a command that is multi-threaded and therefore can utilise multiple CPUs at once. Such programs typically require setting the number of CPU cores to use. Checking that your job is using the number of CPU's expected means your jobs will finish in an optimal time frame. Also, most processing pipelines are a mixture of single threaded and multi-threaded tasks. If your job spends a long time in the single threaded tasks there may be little benefit in requesting a large number of cores for that job. Again, your job could spend a long time in the queue waiting for resources to become available before it's executed, and once those resources are allocated to a job they're not available for anyone else whether they're in use or not.

Monitoring the queues

The *qstat* command shows information about the queue.

```
$ qstat
```

Job id	Name	User	Time Use S	Queue
-----	-----	-----	-----	-----
1220444.tizard1	...build_idx.pbs	jan	0 Q	tizard
1241008.tizard1	Activator	ahao	0 Q	tizard
1241230.tizard1	MARE2DEM	ydidana	0 Q	tizard
1246866.tizard1	zentest	luz	284:06:3 R	tizard
1246867.tizard1	zentest	luz	282:48:5 R	tizard
1246868.tizard1	zentest	luz	283:53:1 R	tizard
1246869.tizard1	zentest	luz	283:59:0 R	tizard
1249511.tizard1	Ce402Pt-1-253	akarayilan	237:01:1 R	tizard
1249514.tizard1	Ce402Pt-1-256	akarayilan	204:09:3 R	tizard
1249527.tizard1	Ce402Pt-1-269	akarayilan	184:29:5 R	tizard
1249549.tizard1	Ce402Pt-1-291	akarayilan	116:43:1 R	tizard
1249554.tizard1	Ce402Pt-1-296	akarayilan	93:25:11 R	tizard
1249557.tizard1	Ce402Pt-1-299	akarayilan	87:43:33 R	tizard

```

1249625.tizard1      ...me.FR07958956 mcorbett      1508:45: R tizard
1249626.tizard1      ...me.FR07958751 mcorbett      1668:27: R tizard
1249627.tizard1      ...me.FR07958750 mcorbett      1185:44: R tizard
1249628.tizard1      ...me.FR07958743 mcorbett      1077:07: R tizard
1249735.tizard1      DirkBoehe          dboehe          21:40:40 R tizard
1249754.tizard1      kartxx_300         akartusinski    17:51:44 R tizard
1249759.tizard1      kartpipe2ph        akartusinski    05:08:15 R tizard
1249764.tizard1      intercept1         tberezowski     108:08:4 R tizard
1249773.tizard1      equilibration      bfielke         17:57:18 R tesla
1249779.tizard1      ...i14_8-8-gIIID  rhudson         08:17:30 R tizard

```

To look at just your own jobs use *qstat -u username*.

To list the available queues and associated memory and maximum walltimes, use the *qstat* command.

```

$ qstat -q
server: tizard1

Queue           Memory CPU Time Walltime Node  Run Que Lm  State
-----
gtx3            23gb  --   50:00:00  --   0   0 --   E R
short           48gb  --   05:00:00  --   0   0 --   E R
gtx             23gb  --   50:00:00  --   0   0 --   E R
tesla           126gb --   100:00:0  --   2   0 --   E R
gpu             94gb  --   50:00:00  --   0   0 --   E R
workshop        48gb  --   05:00:00  --   0   0 --   E R
tizard          --    --   100:00:0  --  374 -219 --   E R

```

376 -219

It is also possible that your job is not able to run and has become blocked. The *showq* command will list any blocked jobs. Jobs are usually only blocked when you've requested resources that will never become available, such as requesting more nodes than are actually present in the cluster, or more memory than is available on a single node.

Check the details of a specific job in more detail using *qstat -f jobid* (just use the numeric part of the JobID).

Monitoring the cluster

The *pbstop* command is a cluster monitoring tool. It provides an overview of the worker nodes in the cluster and displays the current cluster usage, node status, CPU usage and the job queue. You can use it to view the cluster utilisation and available resources.

To show jobs from all users and simply run *pbstop* without additional arguments.

```

$ pbstop

```

To show only your jobs use the `-u` option (replace *username* with your username).

```
$ pbstop -u username
```

You'll need to arrow up and down to scroll across the screen. There are several sections to the output of `pbstop`.

The first line shows that 538/1116 CPUs are presently in use across 21/32 of the worker nodes. The job queue has 164 jobs with 47 currently running.

The second line describes the status of the 32 worker nodes.

```
Usage Totals: 538/1116 Procs, 21/32 Nodes, 47/164 Jobs Running
Node States:  5 down                3 down,offline        18 free                4 job-exclusive        2 offline
```

Node states of “down”, “offline” or “unknown” means the nodes are not presently available in the cluster. “Job - exclusive” and “busy” means all resources of those nodes are being used. “Free” means there are CPU resources available on those nodes.

The second section of `pbstop` is a table which lists all the worker nodes in the left most column (*tizard01*, *tizard02...tizard205*). The numbered columns along the top are the number of CPUs (in blocks of 10) belonging to each node. All Tizard nodes have 48 CPU cores. Each CPU is represented by either a white dot, if the CPU is idle, or a coloured character if the CPU is in use. Each different coloured character refers to a different job which is listed further down the screen.

Only the jobs in the “Running” state will appear on the table.

	1 2 3 4 5 6 7 8 9 0	1 2 3 4 5 6 7 8 9 0	1 2 3 4 5 6 7 8 9 0	1 2 3 4 5 6 7 8 9 0	1 2 3 4 5 6 7 8 9 0	1 2 3 4 5 6 7 8 9 0
tizard01
tizard02	xxxxxxx
tizard03
tizard06
tizard07
tizard08
tizard09
tizard10
tizard13
tizard14
tizard16
tizard17

Monitoring your job

Once your job moves from the “Queued” state to the “Running” state you can see how your job behaves. Just because CPU resources have been allocated to your job doesn’t necessarily mean your script is using all of those resources. Using `pbstop` one can see on which node a job is running. Running the `top` command will the actual CPU and memory usage in real-time.

From *tizard1*, we can use `ssh` to connect to *tizard08* for us and run the `top` command; the output from the `top` command will be sent back to *tizard1* for us to see. Since `top` will display a continuously updating screen, `ssh`

requires the `-t` option. You may also like to add `-u username` (using your own username) after the `-M` option to just list your own processes.

```
$ ssh -t tizard08 top -M
```

(The warning is normal and will occur only the first time you access a node.)

Both of the jobs below have been allocated 24 CPUs, however, one is actually using all 24 cores and the other just a single core. The `%CPU` column shows the CPU usage, where `100%` represents one CPU being fully utilised and `2400%` represents 24 CPU's being fully utilised.

```
top - 16:31:53 up 346 days, 1:58, 1 user, load average: 13.52, 9.94, 9.25
Tasks: 1315 total, 1 running, 1314 sleeping, 0 stopped, 0 zombie
Cpu(s): 49.2%us, 0.1%sy, 0.0%ni, 50.0%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Mem: 132149008k total, 131757136k used, 391872k free, 88188k buffers
Swap: 575635452k total, 54100k used, 575581352k free, 41529908k cached

  PID USER      PR  NI  VIRT  RES  SHR  S %CPU  %MEM    TIME+  COMMAND
 4684 mcorbett  20   0 77308 71g 1092 S 2399.0 5.7 42281.6 bwa
30609 jhack     20   0 14076 2176 876 R 1.6 0.0 0:00.22 top
 196 root      20   0 0 0 0 S 0.3 0.0 23:00.88 events/1
 206 root      20   0 0 0 0 S 0.3 0.0 26:07.39 events/11
 1 root      20  21448 672 464 S 0.0 0.0 0:15.67 init
 2 root      20   0 0 0 0 S 0.0 0.0 0:11.33 kthreadd
 3 root      RT   0 0 0 0 S 0.0 0.0 7:25.17 migration/0
 4 root      20   0 0 0 0 S 0.0 0.0 1:06.00 ksoftirqd/0

top - 16:23:37 up 347 days, 15 min, 1 user, load average: 0.92, 9.53, 10.15
Tasks: 1466 total, 2 running, 1464 sleeping, 0 stopped, 0 zombie
Cpu(s): 2.2%us, 0.1%sy, 0.0%ni, 97.7%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Mem: 132149008k total, 131651764k used, 497244k free, 66616k buffers
Swap: 575635452k total, 35656k used, 575599796k free, 41781716k cached

  PID USER      PR  NI  VIRT  RES  SHR  S %CPU  %MEM    TIME+  COMMAND
37371 mcorbett  20   0 18580 1280 740 R 97.4 0.0 165:58.04 samtools
37372 mcorbett  20   0 77.0g 75g 848 S 8.5 59.9 84:24.70 samtools
37370 mcorbett  20   0 7573m 7.1g 1076 S 3.3 5.7 3322:54 bwa
49380 jhack     20   0 14212 2280 876 R 1.3 0.0 0:00.92 top
 210 root      20   0 0 0 0 S 0.3 0.0 32:13.83 events/15
28249 producer 20   0 361m 11m 2960 S 0.3 0.0 240:06.28 python
 1 root      20  21456 608 412 S 0.0 0.0 0:24.55 init
 2 root      20   0 0 0 0 S 0.0 0.0 0:35.92 kthreadd
 3 root      RT   0 0 0 0 S 0.0 0.0 69:54.07 migration/0
 4 root      20   0 0 0 0 S 0.0 0.0 8:10.83 ksoftirqd/0
```